

In [1]:

```
import cv2
import numpy as np
import matplotlib.pyplot as plt
```

In [2]:

```
import matplotlib.pyplot as plt
import numpy as np
import os
import PIL
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
```

In [3]:

```
import pathlib
data_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz"
data_dir = tf.keras.utils.get_file('flower_photos', origin=data_url, untar=True)
data_dir = pathlib.Path(data_dir)
```

Downloading data from https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz
228813984/228813984 [=====] - 5s 0us/step

In [4]:

```
image_count = len(list(data_dir.glob('*/*.jpg')))
print(image_count)
```

3670

In [5]:

```
roses = list(data_dir.glob('roses/*'))
PIL.Image.open(str(roses[8]))
```

Out[5]:



In [6]:

```
sunflowers = list(data_dir.glob('sunflowers/*'))
```

```
sunflowers = list(data_dir.glob('sunflowers/'))  
PIL.Image.open(str(sunflowers[9]))
```


Out[6]:



In [7]:

```
dandelions = list(data_dir.glob('dandelion/*'))  
PIL.Image.open(str(dandelions[45]))
```

Out[7]:



In [8]:

```
batch_size = 32  
img_height = 180  
img_width = 180
```

In [9]:

```
train_ds = tf.keras.utils.image_dataset_from_directory(  
    data_dir,  
    validation_split=0.2,  
    subset="training",  
    seed=123,  
    image_size=(img_height, img_width),  
    batch_size = batch_size,  
)
```

Found 3670 files belonging to 5 classes.
Using 2936 files for training.

In [10]:

```
test_ds = tf.keras.utils.image_dataset_from_directory(  
    data_dir,  
    validation_split=0.2,  
    subset="validation",  
    seed=123,  
    image_size=(img_height, img_width),  
    batch_size = batch_size,  
)
```

Found 3670 files belonging to 5 classes.
Using 734 files for validation.

In [11]:

```
class_name = train_ds.class_names
print(class_name)
```

```
['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']
```

In [12]:

```
import matplotlib.pyplot as plt
plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(12):
        ax = plt.subplot(3, 4, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_name[labels[i]])
        plt.axis("off")
```

roses



dandelion



tulips



sunflowers




dandelion



roses




dandelion




roses



tulips




dandelion



tulips



tulips



In [13]:

```
for image_batch, label_batch in train_ds:
    print(image_batch.shape)
    print(label_batch.shape)
```

```
(32, 180, 180, 3)
```

```
(32,)
```

```
(32, 180, 180, 3)
```

```
(32,)
```

```
(32, 180, 180, 3)
```

```
(32,)
```

```
(32, 180, 180, 3)
```

```
(32,)
```

```
(32, 180, 180, 3)
```

```
(32,)
```

```
(32, 180, 180, 3)
```

```
(32,)
```

[illegible]

[illegible]

[illegible]

In [14]:

```
AUTOTUNE = tf.data.AUTOTUNE
```

```
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

In [15]:

```
normalization_layer = layers.Rescaling(1./255)
```

In [16]:

```
normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
image_batch, labels_batch = next(iter(normalized_ds))
first_image = image_batch[0]
# Notice the pixel values are now in `[0,1]`.
print(np.min(first_image), np.max(first_image))
```

0.0 1.0

In [17]:

```
num_classes = len(class_name)
model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPool2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPool2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPool2D(),
    layers.Conv2D(128, 3, padding='same', activation='relu'),
    layers.MaxPool2D(),
    layers.Flatten(),
    layers.Dense(256, activation='relu'),
    layers.Dense(num_classes),
])
```

In [18]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
rescaling_1 (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d (MaxPooling2D)	(None, 90, 90, 16)	0
conv2d_1 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 45, 45, 32)	0
conv2d_2 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 22, 22, 64)	0
conv2d_3 (Conv2D)	(None, 22, 22, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 11, 11, 128)	0
flatten (Flatten)	(None, 15488)	0
dense (Dense)	(None, 256)	3965184
dense_1 (Dense)	(None, 5)	1285
=====		
Total params: 4,063,909		
Trainable params: 4,063,909		
Non-trainable params: 0		

In [19]:

```
model.compile(optimizer='adam', loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True), metrics=['accuracy'])
```

In [20]:

```
history = model.fit(train_ds, epochs=10, validation_data=test_ds)
```

```
Epoch 1/10
92/92 [=====] - 80s 863ms/step - loss: 1.3827 - accuracy: 0.3886
- val_loss: 1.0927 - val_accuracy: 0.5477
Epoch 2/10
92/92 [=====] - 80s 866ms/step - loss: 1.0380 - accuracy: 0.5797
- val_loss: 1.0145 - val_accuracy: 0.5831
Epoch 3/10
92/92 [=====] - 80s 865ms/step - loss: 0.9118 - accuracy: 0.6410
- val_loss: 0.8895 - val_accuracy: 0.6376
Epoch 4/10
92/92 [=====] - 79s 856ms/step - loss: 0.7748 - accuracy: 0.6962
- val_loss: 0.8691 - val_accuracy: 0.6621
Epoch 5/10
92/92 [=====] - 80s 870ms/step - loss: 0.6772 - accuracy: 0.7333
- val_loss: 0.9149 - val_accuracy: 0.6703
Epoch 6/10
92/92 [=====] - 80s 870ms/step - loss: 0.5261 - accuracy: 0.8007
- val_loss: 0.9638 - val_accuracy: 0.6608
Epoch 7/10
92/92 [=====] - 78s 853ms/step - loss: 0.4132 - accuracy: 0.8518
- val_loss: 1.1648 - val_accuracy: 0.6267
Epoch 8/10
92/92 [=====] - 80s 869ms/step - loss: 0.2985 - accuracy: 0.8883
```

```
- val_loss: 1.1880 - val_accuracy: 0.6703
Epoch 9/10
92/92 [=====] - 79s 856ms/step - loss: 0.1954 - accuracy: 0.9329
- val_loss: 1.3127 - val_accuracy: 0.6839
Epoch 10/10
92/92 [=====] - 80s 870ms/step - loss: 0.1262 - accuracy: 0.9567
- val_loss: 1.4783 - val_accuracy: 0.6608
```


In [21]:

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs=10
epoch_range = range(epochs)
```

In [22]:

```
plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epoch_range, acc, label='Training Accuracy')
plt.plot(epoch_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epoch_range, loss, label='Training Loss')
plt.plot(epoch_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



In [23]:

```
model.save('Flower.h5')
```

In [24]:

```
flower_url = 'https://storage.googleapis.com/download.tensorflow.org/example_images/592px-Red_sunflower.jpg'
```



```
flower_path = tf.keras.utils.get_file('Red_flower', origin=flower_url )
```

```
Downloading data from https://storage.googleapis.com/download.tensorflow.org/example_images/592px-Red_sunflower.jpg  
117948/117948 [=====] - 0s 0us/step
```

In [25]:

```
PIL.Image.open(flower_path)
```

Out[25]:



In []: